

Kafka

high-throughput, persistent,
multi-reader streams




<http://sna-projects.com/kafka>



- LinkedIn SNA (Search, Network, Analytics)
- Worked on a number of open source projects at LinkedIn (Voldemort, Azkaban, ...)
- Hadoop, data products

Problem

How do you model and process stream data for a large website?



Examples

Tracking and Logging – Who/what/when/
where

Metrics – State of the servers

Queuing – Buffer between online and “nearline”
processing

Change capture – Database updates

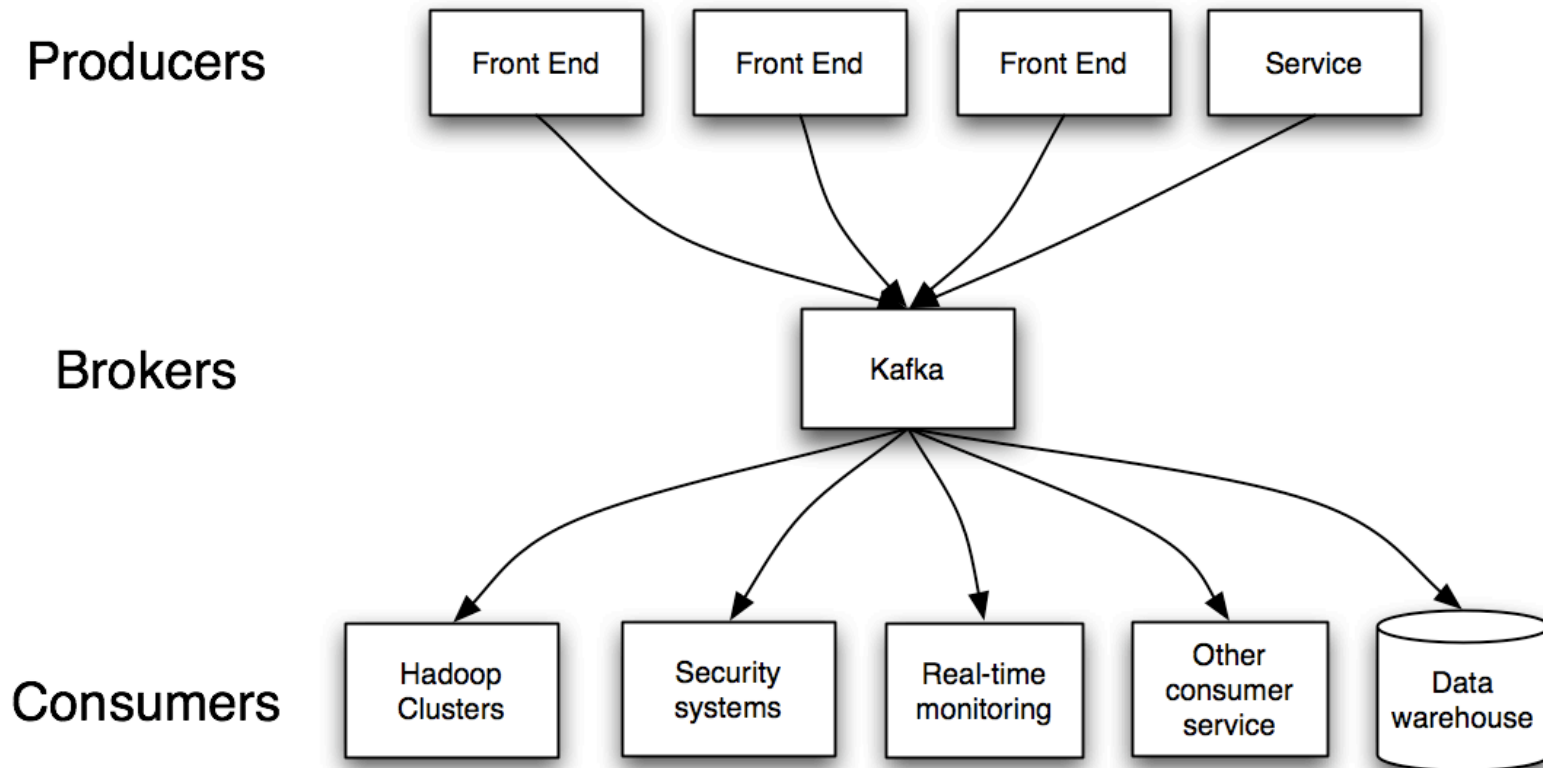
Messaging Examples: numerous JMS brokers,
RabbitMQ, ZeroMQ

The Hard Parts

persistence, scale,
throughput, replication,
semantics, simplicity



Tracking



Tracking Basics

- Example “events” (around 60 types)
 - Search
 - Page view
 - Invitation
 - Impressions
- Avro serialization
- Billions of events
- Hundreds of GBs/day
- Most events have multiple consumers
 - Security services
 - Data warehouse
 - Hadoop
 - News feed
 - Ad hoc

Existing messaging systems


- JMS
 - An API not an implementation
 - Not a very good API
 - Weak or no distribution model
 - High complexity
 - Painful to use
 - Not cross language
- Existing systems seem to perform poorly with large datasets

Ideas


1. Eschew random-access persistent data structures
2. Allow very parallel consumption (e.g. Hadoop)
3. Explicitly distributed
4. Push/Pull not Push/Push

Performance Test


- Two Amazon EC2 large instances
 - Dual core AMD 2.0 GHz
 - 1 7200 rpm SATA drive
 - 8GB memory
- 200 byte messages
- 8 Producer threads
- 1 Consumer thread
- Kafka
 - Flush 10k messages
 - Batch size = 50
- ActiveMQ
 - syncOnWrite = false
 - fileCursor



ActiveMQ vs. Kafka




ActiveMQ vs. Kafka



Performance Summary

- Producer
 - 111,729.6 messages/sec
 - 22.3 MB per sec
- Consumer
 - 193,681.7 messages/sec
 - 38.6 MB per sec
- On on our hardware
 - 50MB/sec produced
 - 90MB/sec consumed




How can we get high
performance with
persistence?

Some tricks

- Disks are fast when used sequentially
 - Single thread linear read/write speed: > 300MB/sec
 - Reads are faster still, when cached
 - Appends are effectively $O(1)$
 - Reads from known offset are effectively $O(1)$
- End-to-end message batching
- Zero-copy network implementation (sendfile)
- Zero copy message processing APIs

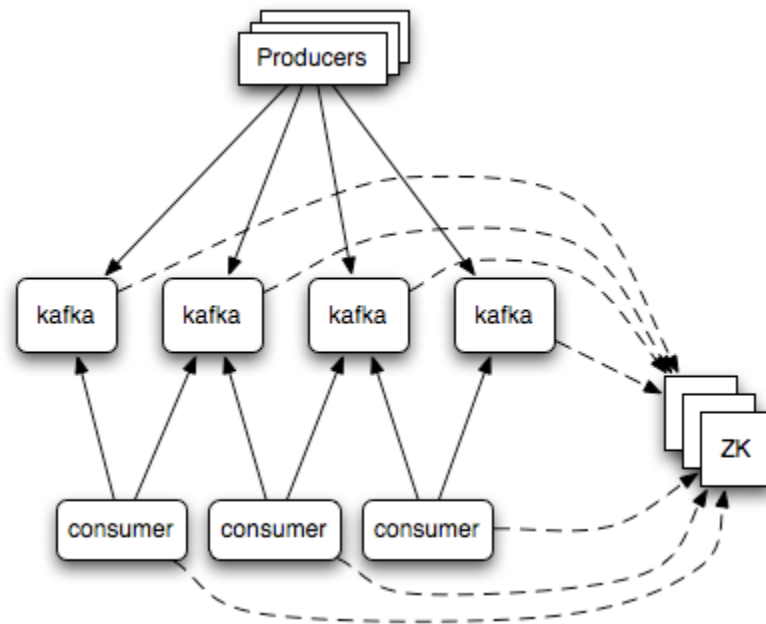
Implementation

- ~5k lines of Scala
- Standalone jar
- NIO socket server
- Zookeeper handles distribution, client state
- Simple protocol
 - Python, Ruby, and PHP clients contributed

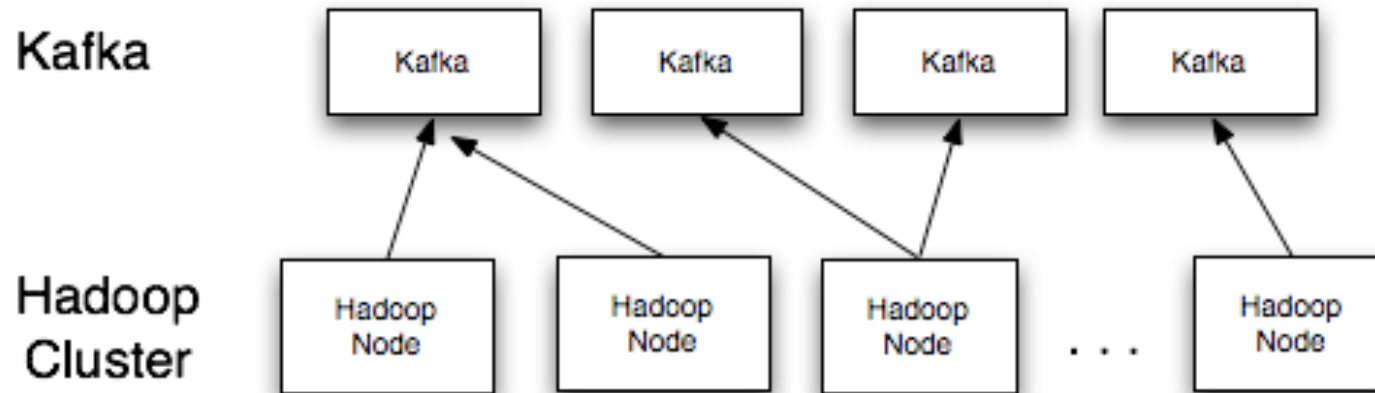


Distribution

- Producer randomly load balanced
- Consumer balances M brokers, N consumers



Hadoop InputFormat



Consumer State

- Data is retained for N days
- Client can calculate next valid offset from any fetch response
- All server APIs are stateless
- Client can reload if necessary



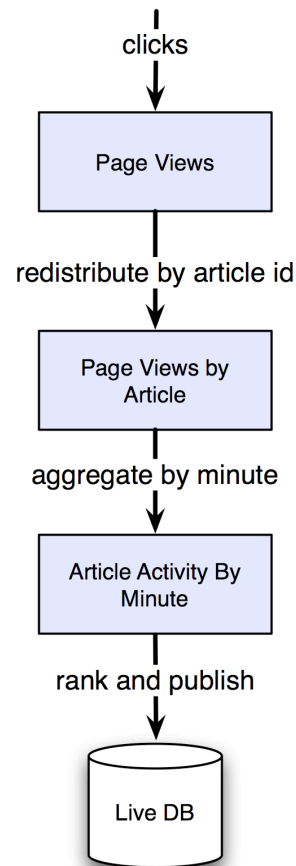
APIs

```
// Sending messages  
client.send("topic", messages)
```

```
// Receiveing messages  
Iterable stream =  
    client.createMessageStreams(...).get("topic").get(0)  
for(message: stream) {  
    // process messages  
}
```

Stream processing (0.06 release)


Data published to
persistent topics,
and redistributed
by primary key
between stages



Rising (Last 24 hours):

1. Reactive oxygen species

Pageviews: 379,200



2. Gerry Ryan

Pageviews: 115,542



3. May Day

Pageviews: 353,067



4. Freddy Krueger

Pageviews: 817,252



5. A Nightmare on Elm Street

Pageviews: 631,561



Also coming soon

- End-to-end block compression
- Contributed php, ruby, python clients
- Hadoop InputFormat, OutputFormat
- Replication



The End

<http://sna-projects.com/kafka>

<https://github.com/kafka-dev/kafka>

kafka-dev@google-groups.com

jay.kreps@gmail.com

